

HUMAN OVERSIGHT IS A CAPABILITY, NOT A CONTROL

What AI Governance Can Learn from Safety Science

Jennifer McCauley • Edward J. Glantz

Penn State University | College of Information Sciences and Technology
ACM AI Summit '26 • Atlanta, GA • Aug 30–Sep 2, 2026



CONNECT

GOVERNANCE CHECKS FOR A PERSON. SAFETY DEPENDS ON A CAPABILITY.

Formal compliance can remain constant while the human capacity to oversee an AI system quietly decays.

1 THE GOVERNANCE ASSUMPTION

“Governance assumes human oversight prevents AI failures.”

Three major frameworks converge on one safeguard:

EU AI ACT : Human oversight

NIST AI RMF : Human review

OECD PRINCIPLES : Human agency

HUMAN OVERSIGHT

Ask whether reviewers retain the time, skill, authority, and diagnostic information to intervene under real operating conditions.

2 THE REFRAMING

Oversight is not a control an organization possesses.

It is a sociotechnical capability it must sustain.

FROM PRESENCE TO PERFORMANCE

BEFORE

Require a human in the loop




AFTER

Measure whether oversight holds

3 OVERSIGHT DRIFT: COMPLIANCE STAYS FLAT. CAPABILITY FALLS.

Safety science predicts that production pressure erodes oversight even when documentation and audit records remain unchanged.



4 WHY THE ASSUMPTION FAILS

1 WORK-AS-IMAGINED ≠ WORK-AS-DONE

Procedures describe ideal work; practitioners adapt under time pressure, workload, and competing demands.

2 NORMALIZATION OF DEVIANCE

Small departures become normal—while compliance records still look unchanged.

3 IRONIES OF AUTOMATION

Automation makes routine work easier but weakens vigilance and expertise when intervention matters.

4 STRETCHED SYSTEMS

Production pressure turns safeguards into resources to be optimized rather than protections to preserve.

5 BUILD GOVERNANCE AROUND CAPABILITY

Four interlocking instruments make oversight observable, maintainable, and resilient:

RELIABILITY METRICS

Test whether reviewers can detect errors, calibrate judgments, and use context.

DRIFT DETECTION

Monitor leading indicators before degraded oversight becomes normalized.

HUMAN-AI SAFETY CASES

Make evidence-backed claims about whether the human role remains safe in context.

MATURITY MODELS

Track where oversight practice is robust, fragile, or merely procedural.

THE TAKEAWAY

A capability that is not measured eventually becomes a checkbox.

Measure oversight reliability. Detect capability drift. Build human-AI safety cases. Mature conditions that let people intervene.

Sources: McCauley & Glantz (2026); EU AI Act; NIST AI RMF 1.0; OECD AI Principles. Poster summary of submitted paper.



PennState
College of Information
Sciences and Technology